

Measuring Secure Coding Practice and Culture: A Finger Pointing at the Moon is not the Moon

Ita Ryan

ADVANCE Centre for Research Training
School of Computer Science and IT
University College Cork
Cork, Ireland
ita.ryan@cs.ucc.ie

Utz Roedig

Connect Research Centre
School of Computer Science and IT
University College Cork
Cork, Ireland
u.roedig@cs.ucc.ie

Klaas-Jan Stol

Lero, the SFI Research Centre for Software
School of Computer Science and IT
University College Cork
Cork, Ireland
k.stol@ucc.ie

Abstract—Software security research has a core problem: it is impossible to prove the security of complex software. A low number of known defects may simply indicate that the software has not been attacked yet, or that successful attacks have not been detected. A high defect count may be the result of white-hat hacker targeting, or of a successful bug bounty program which prevented insecurities from persisting in the wild. This makes it difficult to measure the security of non-trivial software. Researchers instead usually measure effort directed towards ensuring software security. However, different researchers use their own tailored measures, usually devised from industry secure coding guidelines. Not only is there no agreed way to measure effort, there is also no agreement on what effort entails. Qualitative studies emphasise the importance of security culture in an organisation. Where software security practices are introduced solely to ensure compliance with legislative or industry standards, a box-ticking attitude to security may result. The security culture may be weak or non-existent, making it likely that precautions not explicitly mentioned in the standards will be missed. Thus, researchers need both a way to assess software security practice and a way to measure software security culture. To assess security practice, we converted the empirically-established 12 most common software security activities into questions. To assess security culture, we devised a number of questions grounded in prior literature. We ran a secure development survey with both sets of questions, obtaining organic responses from 1,100 software coders in 59 countries. We used proven common activities to assess security practice, and made a first attempt to quantitatively assess aspects of security culture in the broad developer population. Our results show that some coders still work in environments where there is little to no attempt to ensure code security. Security practice and culture do not always correlate, and some organisations with strong secure coding practice have weak secure coding culture. This may lead to problems in defect prevention and sustained software security effort.

Index Terms—Security, secure coding, security compliance, security culture, measuring secure coding

I. INTRODUCTION

Software security was defined in 2004 as ‘the idea of engineering software so that it continues to function correctly under malicious attack’ [1]. Since then, our world has become increasingly connected, and the hacker community has morphed into a global criminal industry. Insecure software is now routinely exploited for ransomware, cybercrime and cyber-

espionage purposes, so that software security is an increasingly urgent requirement [2].

Academia has not reached consensus on a way to measure the level of secure coding in an organisation or open source team. Ethnographic studies provide rich qualitative data, but by their nature [3] do not produce quantitative measurements or generalisable conclusions. Most researchers attempting quantitative secure coding research synthesise a number of practices from industry methodologies and assess participants’ use of them. These synthesised secure coding practice lists tend to be used only by their authors. They are not empirically evaluated, and if fine-grained may be too rigid to be re-used over time, as software security is a quickly-evolving field.

To explore these issues, we conducted a large-scale survey of software developers. In a recent analysis of a large trove of industry secure-coding data, twelve security activities were identified by Weir et al. [4] as being those most commonly adopted in secure coding initiatives by large organisations. Given their empirical backing, we adopted these 12 Common Activities (CAs) as our core measurements for secure software coding practice.

To measure security culture we devised a list of questions grounded firmly in prior literature.

In this paper we seek to address the following research questions:

Research Question 1: *What are the secure-coding characteristics of our sample group?*

Research Question 2: *What are the security culture characteristics of our sample group?*

Research Question 3: *Do secure coding practice and culture correlate, and if not, what lessons can we learn to help support the development of secure coding?*

This paper makes the following contributions: first, we present the ‘CA Score’, a lightweight instrument for assessment of the level of software security practice in a work environment. Second, we provide an overview of the current state of security practice based on our survey results. Third, we introduce a group of questions designed to probe security culture.

Fourth, we assess security culture in our participants' working environments. Finally, we explore the interplay between security practice and security culture.

The paper is organised as follows: background and related work are discussed in Section II, method is discussed in Section III, an overview of the sample is given in Section IV, results are given in Section V. Section VI contains the discussion and threats to validity, and concludes the paper.

II. BACKGROUND AND RELATED WORK

A. Measuring Secure Development

While assessing the security of software developed in controlled studies is feasible and repeatable [5]–[7], measuring secure development in organisations and open source environments is an ongoing issue in academia. As far back as 2012, Jaatun [8] discussed whether software security can be measured at all. Counts of known security defects for publicly available applications could be used; however, some applications are subject to considerably more analysis and attack than others. An absence of known vulnerabilities could simply indicate that the application has not been closely studied. An example of an attempt to demonstrate a correlation between use of assurance techniques and decreased levels of security issues in code comes from Weir et al. [9]. They surveyed 335 Android app developers on their use of security assurance techniques, downloaded a free app from each developer, and used automated analysis to detect instances of three categories of security issue in the apps. Surprisingly, app analysis showed no increased security for apps whose creators claimed to use assurance techniques, and more cryptographic security issues for apps whose creators had access to security experts. The authors concluded that issues with analysis tools and missing cryptography may have skewed the analysis. This study illustrates the complexity and difficulty inherent in determining whether secure coding efforts are correctly reported, and whether they result in secure code outcomes.

BSIMM reports [10] on contemporary software security activities are compiled semi-annually by Synopsys, who interview contributing organisations about their security practices. For example, in 2021 BSIMM 12 collated data from 128 organisations. The contributing organisations are usually large and have active software security strategies. Jaatun [8] discussed adopting a BSIMM-like analysis method for academic studies. Variations on this approach of measuring activities from BSIMM and other industry practices and guidelines such as OWASP's Software Assurance Maturity Model (SAMM) [11], Microsoft Security Development Lifecycle (MS-SDL) [12] and Software Assurance Forum for Excellence in Code (SAFECode) [13] have been widely used by the research community [14]–[19]. BSIMM currently includes a total of 122 activities, with correspondingly large numbers of similar activities in the other methods. This is cumbersome, therefore a synthesised subset is usually generated. Each research team measures different activities. New groups of activities to measure are continually defined. None are empirically evaluated.

Research teams rarely reuse the groups defined by previous teams, making results difficult to compare.

Similarly, Morrison et al. [20] found that 85% of 324 unique security metrics had been proposed and evaluated solely by their authors. They concluded that there is no convergence as yet on an accepted set of metrics in the software life cycle security metrics field.

B. Security Culture

It might be argued that engaging in security practices would naturally entail a good security culture. However, many observers have noted that organisations may follow recommended practices (perhaps for compliance) but have a box-ticking attitude to security. Although mentioned in several papers [21], [22], and implied in others [23]–[25], this phenomenon and its implications are rarely explored. It may lead to the implementation of security activities purely to demonstrate compliance with regulations or standards such as the payment card industry's DSS [26]. Such standards may have a narrow focus [25], and security investment which focuses solely on compliance is not sufficient to ensure secure software [27]. The significance of the organisation's security posture in the area of secure coding is widely acknowledged in academia [28]–[30]. Assal et al. [31] found that most deterrents to developers coding securely relate to the absence of clear secure-coding plans, resources and priorities in the developers' working environments. Concluding that it is important for companies to build a security culture and offer learning opportunities for secure coding, they suggested that a research focus on problems in the organisational approach to security would increase understanding of software security deterrents.

Recent ethnographic studies have provided useful insights into the software security posture in organisations. Lopez et al. [32] immersed themselves in the software development unit of a large company, examining how non-experts treat security during their daily tasks and focusing primarily on developer security motivation. They concluded that the organisation had put procedures and activities in place to ensure software security, and that the developers accepted these and attempted to implement them in good faith. This resulted in '*mostly secure*' software. By contrast, Morales et al. [24] published a devastating paper on a well-funded project using Agile and implementing DevSecOps without rigorous adherence to DevSecOps principles, and the resulting security implications. It is a valuable reminder that no development method can compensate for bad management. Broken pipelines, adversarial subcontractor relationships, inadequate definitions of done, poor test resources, a failure to emphasise security and more all combined to produce a product with poor security assurance. On paper, this dysfunctional project adopted appropriate security practices. So where did it go wrong?

The differentiating factor between these two organisations is security culture, defined by Haney et al. [22] as '*a subculture of an organization in which security becomes a natural aspect in the daily activities of every employee.*' Haney et al. found that all of the cryptography-focused organisations they studied

placed a high value on a strong security culture. Tuladhar et al. [33], having worked embedded in an organisation adopting secure coding practices, concluded that key security culture enablers were upper management setting security as a goal, and the team applying security knowledge in context. In a longitudinal study, Tøndel et al. [21] noted the adverse effect of managerial lack of interest in security. They also observed that developer security work was often invisible to others and therefore did not help build security culture.

Managerial attitude affects the question of how individual developers' security enthusiasm is regarded within their work environment, an important aspect of security culture that has not been much explored by the research community. Jaatun et al. [14] noted organisations' dependence on individual developers' enthusiasm. Tahaei et al. [34] examined the experience of privacy '*champions*' in software teams, finding that they play an important advocacy role. Ryan et al. [35] identified a '*hero*' software security archetype; a coder struggling to introduce secure coding practices in a security-hostile environment. Such environments are difficult to study, since most security experts work in roles where security is already a focus [36]. Security interventions, such as those introduced to organisations by Weir et al. to attempt to empower developers to code securely [37], require prior organisational agreement. Thus a level of management support is assured.

Developers' tool adoption behaviour may give useful insights into the work environment with which they must contend. Papers on tool adoption have focused on social influences [38]. Although Witschey et al. [39] observed that many developers seek out information on security tools when needing to write secure code, they did not examine how developers can introduce such tools into their working environment, and what constraints they encounter when attempting to do so. An ethnographic study by Palombo et al. [40] shed some light on how heroes can succeed in improving security in a '*security inattentive*' [18] environment. Two researchers used and advocated a '*co-creation*' model for secure development in which they added security checks seamlessly, with no developer friction for other team members. While this required considerable investigation and work, other team members then adopted the checks, possibly because of their frictionless nature.

The interplay between developer and work environment is the key to security culture. Therefore, ways to measure it are of interest. One suggested measurement tool is the Secure Software Development Self-Efficacy Scale (SSD-SES) [19], although a developer's self-efficacy can arise from previous employment or personal motivation rather than their current work environment.

C. Organisational Climate Theory

Academic work on organisational culture originated from anthropology and considers shared myths, beliefs and ritual [41]. The term '*security culture*' as used in software security literature is adopted from industry [22], and may have more in common with organisational climate theory, which provides a way to consider and evaluate the organisational '*climate*' for

differing organisational '*facets*.' For example, an organisation may demonstrate a consistent commitment to the software security facet by halting releases when there is a security concern. Arizon-Peretz et al. [42] applied this theory to software security, asking developers about privacy and security activities in their organisation and interpreting the answers via organisational climate theory. They categorised activities into seven climate themes which provide cues to developers as to the relative importance of security and privacy within the organisation, and what is expected of them in these areas. They found that the cues that developers get can be confusing and do not clearly indicate that security and privacy should be prioritised in their daily work. They proposed that climate theory can be used to change this and to promote security and privacy by design within organisations. In a follow-up paper [43], they measured organisational climate for the software security facet using questions derived from the seven climate themes and tailored to a multinational organisation. They found an interplay between security self-efficacy, proactive security behaviour and a positive organisational climate for software security within the organisation.

III. METHOD

We conducted a large-scale cross-sectional survey. In this section we discuss the design of the questionnaire, data collection procedures, data screening procedures, and data analysis procedures. Data and scripts are available online [44].

Surveys that constrain respondents' answers to specific values can impose forced choice bias, which is inappropriate in a context where increased understanding of context is sought [45]. Therefore we allowed free text in many questions, and asked a question towards the end looking for comments on any aspect of the survey or of the respondent's secure software development experience. The free text aspect of the survey provided us with many interesting insights. As the use of '*sic*' can be seen as condescending, we do not use it when quoting respondents. Any quote from respondents is presented exactly as entered in the questionnaire.

A. Survey Questions

Initial questions concerned demographics and participants' personal relationship with secure coding. These questions are not the focus of this paper and most are not discussed. Some standard secure coding questions were asked; answers are discussed in Section IV.

1) *Screening Questions*: One concern in sample research is that the sample is representative. We were only interested in responses from people who were coding frequently at the time of answering the questionnaire, so we began by asking respondents '*Do you write computer code frequently, either professionally or open source?*' Those who answered '*No*' were excluded from further participation.

Researchers running secure software development surveys have encountered issues with respondent quality when offering payment for completion. Witschey et al. [46] cleaned 313 suspiciously-speedy participants from their survey, leaving

only 61 usable responses. Danilova et al. [47] found that of 129 respondents sourced from Qualtrics, 96 did not pass programming tests. As a result, Danilova et al. [48] assessed a number of screening questions to filter out non-developers from paid studies. The Danilova paper recommended two programming comprehension questions as the best screening choices in surveys where a time penalty is acceptable. These two questions have since been used successfully, for example by Kaur et al. [49]. We adopted the two questions, adapting the wording slightly since the answers documented by Danilova et al. are in the public domain. We had some interesting findings on them, discussed in section III-C.

Since our questionnaire was relatively long with 62 questions, we wanted to ensure that participants were paying attention all the way through. Rather than ask complex attention questions with opposing answers, we borrowed a simple strategy from Murphy-Hill et al. [50], simply stating for example *'This is an attention question. Please select a little.'* We had two such questions, designed to blend with adjacent questions so that participants who were simply clicking blindly would not notice them. This strategy caused some amusement; for example, one participant noted as a comment at the end of the survey *'The attention questions are funny ^^'*

2) *Measuring Security Practice*: We asked our participants whether they were aware of each of the 12 CAs in their working environments. These questions can be found in Table III. More detail on the rationale for these questions can be found in section V-A.

3) *Measuring Security Culture*: Security culture questions were based on an extensive literature review of software security papers since 2016. These questions can be found in Table IV. More detail on how these questions were devised can be found in section V-B.

4) *Pilot Testing*: The questionnaire was pilot tested in two phases. First, we asked for criticism and feedback from a developer with over two decades of experience. After making recommended changes, we pilot tested again with the same developer and two other highly experienced developers. Small adjustments suggested by these testers were made before launching the questionnaire. Their responses were not included in reported results.

B. Data Collection

The survey was publicised via personal contacts, a conference talk by the first author, and social media platforms such as LinkedIn, Facebook, and Twitter. We found that it was possible to generate significant interest by asking for retweets on Twitter, posting to Facebook development groups, and promoting the survey on LinkedIn. The survey was open for just over 3 months, to the end of January 2022. This long window provided time to publicise the survey, and helped build momentum. At closing time we had received 1,100 responses.

Given the considerable interest we succeeded in generating in our survey, we did not need to offer financial incentives or use platforms such as Prolific to source participants. This removed the danger of non-developers participating for financial gain.

It remained important to carefully screen responses to ensure high quality analysis; we discuss the screening process next.

C. Screening Process

As only the questions on programming expertise were mandatory, and some people think more quickly than others, we did not remove responses that were completed quickly.

Fifty-two respondents answered 'No' to the initial screening question on frequent coding and were brought directly to the end of the survey, leaving 1,048 respondents. We then proceeded to the two mandatory programming expertise questions. We had intended to reject all participants who got either of the programming questions wrong. However, a tranche of six or seven survey entries was obtained immediately after the first author gave a software security talk to an audience of SQL programmers. These respondents included SQL in their list of programming languages, and some got the second programming question wrong. Upon reviewing the questions, we realised that these included pseudo-code that would be typical of C++ or Java interview questions but is meaningless in SQL. The questions' limitations were confirmed later, when a survey participant communicated with the first author, commenting that Python programmers could have difficulty with these questions. Fifty-five respondents answered one or both of the programming screening questions incorrectly. Similarly, we had planned to remove any respondent from the data set who got an attention question wrong. This seemed like an uncontroversial position, but one participant, answering Question 60 which gave an opportunity to comment on the survey, responded *'I didint understand q56. Attention? No idea what you are wanting from this poor question.'* Thirty-seven respondents failed one or both of the attention questions.

Having considered the feedback on our screening questions, we decided to consider two groups of respondents during analysis. Group 1 (*'all valid participants'*) includes all respondents who confirmed that they write computer code frequently ($n=1,048$). Group 2 (*'all correct participants'*) additionally omits respondents who got programming and/or attention questions wrong (but retains the respondent who explicitly told us that they didn't understand the attention questions), giving a total of 962. Statistical results reported in the paper are based on Group 2; however, we also ran tests on Group 1. We found no significant differences between the two groups.

D. Ethics

Our questionnaire involved research on human subjects, and therefore we obtained approval for the questionnaire from our institution's Social Research Ethics Committee. Participants were advised that taking the questionnaire was not obligatory, and were asked to consent to taking the questionnaire. All submissions were anonymous.

E. Data Analysis

Statistical analysis was done using R [51]. Positive answers to the 12 CA questions were summed to create a CA score between 0 and 12, representing the secure coding level in




Fig. 1. Age distribution of survey respondents

TABLE I
GENDER DISTRIBUTION OF SURVEY RESPONDENTS

Gender	Frequency	Percent
Man	852	81
Woman	94	9
Non-binary	63	6
Prefer not to say	34	3
Other	4	0.4

an environment. The Pearson correlation coefficient was used to assess correlations between this score and security culture indicators, measuring correlation of increasing scores with increasing levels of these indicators [52]. Participants were asked to choose ‘*Not Applicable*’ where relevant; for example, if they worked alone and the practice required co-operation. For the Common Activities, participants were also provided with an ‘*I don’t know*’ option. Levels of ‘*I don’t know*’ responses varied from 5.93% (CAQ1) to 22.6% (CAQ12). Since levels were high for some questions, we ran correlations excluding NA and ‘*I don’t know*’ responses. We did not find any differences in these results that would affect the paper’s conclusions. Therefore, the correlations given in the results section do not omit ‘*Not Applicable*’ or ‘*I don’t know*’ answers. Qualitative data from open-ended survey questions was used to provide context for our quantitative results.

IV. SAMPLE DESCRIPTION

A. Demographics

Respondents ranged in age from 18 to over 65 (see Fig. 1). We asked respondents what gender they most identified with. The majority of respondents answered ‘*Man*.’ The numbers seemed to correspond with approximate gender balance numbers in the industry (see Table I). We asked developers what country they live in, allowing them to add their country if it was missing. We obtained responses from 59 different countries across all continents except Antarctica. The highest numbers of respondents came from the United States (453), United Kingdom (110) and Germany (92).

B. Programming Languages and Tools

We provided developers with a list of programming languages and other technologies and asked which they used frequently, also allowing free text entries. We ended with 73 programming languages, of which JavaScript (454), Python (447) and HTML/CSS (332) were the most popular. Some single-user languages, such as CHICKENscheme and MoonScript, were

entirely new to us. We asked about other technologies; 129 were entered. Git (966) was by far the most popular, followed by Docker (470), PostgreSQL (349) and Amazon Web Services (AWS) (336). Sixty-eight unusual technologies appeared only once, including Cassandra, ScyllaDB, and Godot, which is apparently ‘*The game engine you waited for*.’

The field of tools used to enhance code security is rapidly evolving. We wanted to get a sense of which are in general use. Asked which security tools they use, developers entered over 100 different tools. SonarQube (196) and Clang Analyzer (175) were by far the most popular, but the large number of tools indicates little convergence in the tools market.

C. Adoption of Secure Coding Standards

We asked developers what secure coding initiatives or standards they are aware of in their organisations or teams. We provided a list of well-known initiatives, and allowed developers to add their own. By far the most common response was ‘*None*’ (304), followed by ‘*Not applicable*’ (212). PCI-DSS, used in the payments industry, followed at 78, with FIPS 140-2 (43), Common Criteria (37) and US-CERT’s Top 10 Secure Coding Practices (33) also fairly common. Although included in the predefined list, industry secure-coding standards referenced in academia fared badly; the BSIMM got only three mentions, OWASP’s SAMM two, and SAFECODE five, and of the commonly-referenced general approaches only MS-SDL hit double digits at 29.

The free text section provided further insights. MISRA, standards set by the Motor Industry Software Reliability Association, appeared seven times. Mention was also made of government standards from France, Germany, the US and the UK. We had inadvertently omitted an ‘*I don’t know*’ option; 18 people added this in free text. We received other thoughtful entries such as ‘*Standards are artificial, we try to focus on actual security (incidentally we probably apply some of them)*’ and some less thoughtful ones: ‘*Yes there is some I don’t care*.’

D. Adoption of Development Methods

We asked participants about the development methods in their organisation, providing 15 common options (allowing selection of multiple options), and a free text option. We found that Agile (650), DevOps (445), Scrum (358) and Kanban (296) were the most common responses. Some of the 28 participants who used the free text option were positive, with comments such as ‘*Good practices but no dogmatic application of any of these*.’ The majority of free-text respondents delighted in letting off steam, for example: ‘*...They claim it’s scrum / agile. It is not, it’s waterfall with extra meetings*,’ ‘*Fake cargo cult agile*,’ and ‘*RDD (Resume-Driven Development)*.’

E. Secure Coding Policy

Asked whether their organisation or team has a written secure coding policy, less than a fifth of our participants (184) answered ‘*Yes*.’ Other preset options were ‘*No*’ (537), ‘*I’m not sure*’ (238) and ‘*Not applicable*’ (65). Testament to the positive impact of security surveys, one free text respondent wrote: ‘*I’m going to write one today*.’

F. Security Priorities

We asked people what aspects of software security are important. Data protection (978), Preventing vulnerabilities (954), Customer privacy (877) and Customer confidentiality (829) were widely chosen. Most other aspects were selected by at least 30% of respondents. Only two people chose *‘I do not think that software security is important.’* There were 49 free text entries comprising a wide range of priorities and experience, usefully summed up by the single entry *‘It Depends™.’*

G. Training

We asked participants whether they had ever been offered security or privacy training, and the majority answered *‘No’* (see Table II). We then asked for training details, and received over 300 individual free text entries. They ranged from negative comments such as *‘Boring’* to more positive feedback: *‘Nothing formal. It happens ad hoc (usually directed by me) as a part of code review.’* A reminder of the reluctance of individuals who take security seriously to part with confidential information, several answers were inscrutable, e.g. *‘Confidential,’* *‘Information refused,’* and *‘Internal training.’*

TABLE II
RESPONDENT HAS BEEN OFFERED SECURITY OR PRIVACY TRAINING

Response	Frequency	Percent
Yes	416	40
No	525	51
Not applicable	98	9

V. RESULTS

One of our survey participants, asked about the importance of security activities, responded that *‘External checks are a “finger not the moon” problem, but a useful diagnostic.’* This is a reference to a Buddhist saying on dogma, *‘A finger pointing at the moon is not the moon,’* suggesting that it is easy to confuse descriptions of a thing for the thing itself. Given the difficulty in ascertaining whether software is secure, and the comparative simplicity of evaluating lists of software security activities, it is easy to confuse the activities with the goal. Evaluating security culture alongside security activities may get us a little closer to secure coding reality.

A. Research Question 1: What are the secure-coding characteristics of our sample group?

As discussed in Section II, there is no agreed measurement of software security practice in academia. Different groups of academics roll their own, based on BSIMM, SAMM, MS-SDL and other practices. Measurement tools are thus not empirically validated, are not easily comparable between studies, and tend not to be reused across studies.

Our empirical approach was motivated by Weir et al. [4], who analysed the BSIMM assessments dataset to study how security activities were introduced in organisational settings over a period of 12 years. They found that there were 12 activities that were adopted *‘most often, together and first’* by

a majority of organisations, with at least half of them found in 92% of the assessments. There was *‘a marked jump’* of 18 percentage points between the frequency of use of the other developer activities studied, of which the highest frequency was 47% (i.e. less than half), and these 12 activities, which had frequencies from 65% (CAQ5) to 89% (CAQ10). We judged that determining the presence or absence of these activities in an environment would give a contemporary objective assessment of the environment’s software security practice. (Note that although named and categorised by the BSIMM investigators, these activities are common in industry and their adoption does not depend on familiarity with the BSIMM).

We asked respondents about the presence of these 12 known most Common Activities (CAs) in their environment. See Table III for the exact text of each of the 12 questions. We evaluated the number of the CAs that they identified as in use in their working environments (see Fig. 2). We assigned a score of 0-12 to work environments by counting the number of CAs they were undertaking from the 12 we asked about in the survey. We call this the *‘CA Score.’* This score is empirically justifiable, practical, and repeatable, and thus fills a gap in secure coding measurement.

Percentages of respondents choosing *‘True’* to the CAs are in Table III. A full breakdown of answers can be found in the online supplemental material. While a small number of participants stated that all 12 CAs were used, most used fewer, and quite a few participants indicated that none or very few of the activities were undertaken in their work context.

B. Research Question 2: What are the security culture characteristics of our sample group?

We asked a series of questions about security culture which can be found in Table IV. The questions were based on themes related to security culture mentioned in multiple papers in the existing literature but which appeared to be largely unexplored. The rationale for each question is given in the discussion of the questions below.

1) *Support for Secure Coding (SCQ1):* Support from the organisation has been found to be a significant aspect of security culture in multiple studies [18], [22], [31], [33]. We asked to what extent participants felt supported to code securely (SCQ1). Fig. 3 presents the answers to this question. As can be observed, answers are fairly evenly divided between those who feel supported, those who do not feel supported, and neutral




Fig. 2. Number of security Common Activities (CAs) undertaken in respondents’ coding environments.

TABLE III
QUESTIONS ON TWELVE COMMON ACTIVITIES

No.	True	Question
CAQ1	58.5%	Static analysis tools are brought into the code review process to make the review more efficient and consistent.
CAQ2	41.1%	Compliance constraints are translated into software requirements for individual projects and are communicated to the engineering teams.
CAQ3	36.2%	QA efforts go beyond functional testing to perform basic adversarial tests and probe simple edge cases and boundary conditions, with no particular attacker skills required.
CAQ4	31.3%	QA targets declarative security mechanisms with tests derived from requirements and security features. A test could try to access administrative functionality as an unprivileged user, for example, or verify that a user account becomes locked after some number of failed authentication attempts.
CAQ5	28.7%	Penetration test tools are used internally.
CAQ6	62.1%	Emergency codebase response can be done. The organisation or team can make quick code and configuration changes when software (e.g., application, API, microservice, CAQ infrastructure) is under attack.
CAQ7	67.8%	Defects found in operations are entered into established defect management systems and tracked through the fix process.
CAQ8	77.5%	Bugs found in operations monitoring are fed back to development, and may change developer behaviour. For example, viewing production logs may reveal a need for increased logging.
CAQ9	36.5%	Penetration testing results are fed back to engineering through established defect management or mitigation channels, with development and operations responding via a defect management and release process.
CAQ10	65.8%	Host and network security basics are in place across any data centers and networks and remain in place during new releases.
CAQ11	34.9%	Security-aware reviewers identify the security features in an application and its deployment configuration (authentication, access control, use of cryptography, etc.), and then inspect the design and runtime parameters for problems that would cause these features to fail at their purpose or otherwise prove insufficient.
CAQ12	31.7%	External penetration testers are used to identify security problems.

responses. Answers to this question had a weak to moderate correlation of .46 with CA scores ($p < .001$).




Fig. 3. SCQ1: How much support do you feel that you get from your employer or open source team to code securely? (n=877).

2) *Raising Security Concerns (SCQ2)*: We asked participants whether, if they had a security concern at work, they would raise that concern. This question is part of our investigation into how security-motivated developers can influence their

TABLE IV
SECURITY CULTURE QUESTIONS

ID	Question
SCQ1	How much support do you feel that you get from your employer or open source team to code securely?
SCQ2	If you had security concerns about your current project, how likely would you be to raise them with someone?
SCQ3	Security tools are tools that help to ensure source code is free from security vulnerabilities. For example, they may analyse code for known issues or check configurations for problems. Some examples are Coverity, CodeSonar and SonarQube. Some security tools like SonarLint integrate with the IDE. Are there security tools available in your environment?
SCQ4	If you saw a free tool that you felt would help you to code more securely, how likely would you be to install and use it without asking anyone for permission?
SCQ5	If you needed permission to use the tool, how likely would you be to ask for permission?
SCQ6	If you asked, how likely do you think it is that you would get permission?
SCQ7	If you asked, how likely do you think it is that you would get funding for a paid tool?
SCQ8	Have you ever heard people say that there is a software security culture in your working environment?
SCQ9	Would you agree that there is a security culture in your working environment?
SCQ10	How highly do you think your team prioritises software security?
SCQ11	How often is software security mentioned in team communications?
SCQ12	Roughly how much time do you spend on software security in an average week? This would include any tasks aimed at making or keeping a product secure, such as fixing vulnerabilities that could cause a breach, keeping third-party components updated and assessing code for security issues.

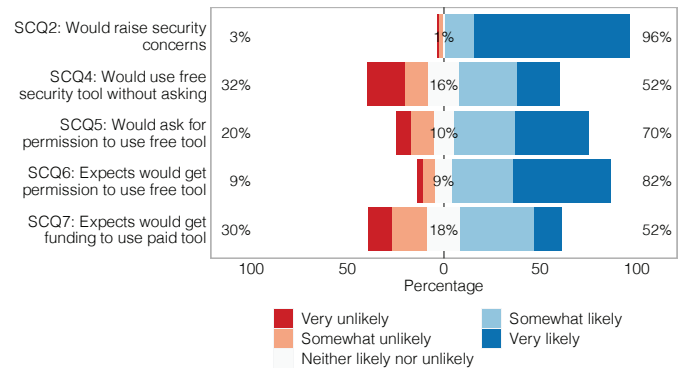


Fig. 4. Results from questions SCQ2 (n=959), SCQ4 (n=959), SCQ5 (n=858), SCQ6 (n=860) and SCQ7 (n=891). Questions concern the reactions respondents anticipate if they were proactive about security.

working environment. Fig. 4 shows that the vast majority of participants would be somewhat or very likely to raise the concern. Those who would not may work in organisations or teams where raising security concerns is actively discouraged, a phenomenon that has previously been observed [18]. This question can be a useful way to detect coding environments

that are extremely unfavourable to security. Answers to this question had a correlation of .25 with CA scores, with $p < .001$.

3) *Whether there are Security Tools in the Working Environment (SCQ3)*: Security tool use is a fundamental aspect of secure coding, mentioned in most of the relevant literature on the subject [53]–[55]. While it may also be used to assess security practice, tool use is an effective proxy for the cultural value of expending effort, time and money on security. Therefore, we asked whether security tools are present in the developer's environment. See Fig. 5, which shows that 33% of respondents had no security tools present in their working environment. Answers to this question had a low correlation of $r = .36$ with CA scores ($p < .001$).




Fig. 5. SCQ3: Are there security tools available in your environment? (n=932).

4) *Introducing Free Tools (SCQ4)*: Previous research on reasons why developers adopt security tools focused on their motivations and inspiration to do so. In this survey we wanted to explore the level of support for finding and integrating such tools that developers are likely to find within their environment. Xiao et al. [38] noted that several participants in their widely-cited interview study on how developers adopt security tools were self-motivated, seeking out tools because they needed them, and finding them online. Those developers would then need to introduce the tools to their work environment. This process of introduction was not a focus of the Xiao et al. or related Witschey et al. [39] paper, and remained unexplored. In a follow-up survey by Witschey et al. [46], the authors expressed surprise that although statements expressing 'security concern' were predictors for security tool use, they had weak effects. We suspected that this could relate to a difficulty for developers in introducing security tools. In security-conscious environments, there could be a prolonged approval process or outright prohibition on new practices and tools. In security-hostile environments, they could be seen as a waste of time. The interaction of security-motivated developers with their environment influences their security activities [47], [56], yet the degree of support or resistance they are likely to encounter has not been investigated. Therefore we asked several questions in our survey about attempting to introduce new security tools. Questions SCQ4-SCQ7 explore how developers would go about introducing such tools, and what reaction they anticipate they would get within their environment. Developers' subjective expectations are important, since a perception that tools would be rejected could result in their not seeking out or proposing them [42]. An organisational climate that is positive for software security has been shown to be positively related to proactive developer behaviour [43].

SCQ4 asked whether the respondent would use a free security tool without asking for permission. See Fig. 4 for the results;

52% of those surveyed would use such a tool, while 32% would not. Note that while it might be expected that security-aware organisations would encourage experimentation with security tools, use of a free tool could itself introduce security vulnerabilities. Ironically, many security tools run with elevated privileges and can be sources of risk. Answers to this question did not correlate at all with CA scores ($r = -0.08$, $p = .009$).

5) *Asking for Tool Permissions (SCQ5-SCQ7)*: In SCQ5, we asked participants how likely it was that they would ask for permission to use a free tool (see Fig. 4). Those selecting 'Not Applicable' were discounted. 70% of respondents would ask, 10% were neutral and 20% were unlikely or very unlikely to ask. Thus, 20% of respondents apparently have a particularly low interest in security, particularly low expectations from their management team, or both. (Answers to this question did not correlate with CA scores ($r = .15$, $p < .001$)).

However, SCQ6 indicated that a large 82% of respondents thought that it was somewhat or very likely that, if they did ask, they would get permission for a **free** tool. The figure for funding for a **paid** tool (SCQ7) was much lower at 52%. It can be argued that paid security tools are often better supported than free tools, so that they should be preferred. The only negative difference between a free and paid tool is financial outlay. The contrast between the answers to these two questions may indicate a poor security culture in the relevant participants' organisations. Haney et al. [22] included spending money on security in their list of attributes needed for a positive security culture in a development environment. The signals and cues received by employees influence their understanding of the real priorities in their work environments [31], [42]. In this case, participants have concluded that security is not a management priority.

Answers to SCQ6 did not correlate with CA scores ($r = .11$, $p = .001$), but SCQ7 answers showed a weak correlation of .27 ($p < .001$), indicating that paid tools are slightly more likely to be considered if other security precautions are in place.

6) *Perceptions of Security Culture (SCQ8-SCQ9)*: Although the importance of security culture is emphasised in numerous papers on organisations and teams with secure coding processes [21], [23], [33], [57], [58], there can be a disconnect between management and developers when it comes to security [23]. Some researchers have found that developers can be cynical about management drives for a security culture [58], perhaps perceiving them as lip service. Therefore we first asked whether the developer is aware of a claim to security culture in the environment, and then asked whether the developer would agree with this claim. Asked whether they had heard it said that there was a security culture in their working environment (SCQ8), only 261 respondents said 'Yes'; 633 selected 'No,' with 133 selecting 'Not applicable.' We allowed free text contributions; one was: 'no real people say stuff like that,' a reminder that the term 'security culture' can seem like just more management speak or propaganda [58] if its use is not accompanied by concrete steps towards prioritising security. Another contribution was 'We're a small firm, shipping working software at all is the priority.' Answers to this question

correlated weakly with CA scores, ($r = .40$, $p < .001$).

The following question (SCQ9) was whether participants would agree that there was a security culture in their work environment. As can be seen in Fig. 6, 48% of participants agreed or strongly agreed with this statement. 28% disagreed or strongly disagreed, while 24% were neutral. This question correlated with the CA score, with a weak to moderate correlation ($r = .41$, $p < .001$).




Fig. 6. SCQ9: Would you agree that there is a security culture in your working environment? (n=946).

Of particular interest is whether there is a strong correlation between how security culture is portrayed in a work environment (SCQ8) and how respondents themselves view the environment's security culture (SCQ9). In fact, there is only a moderate correlation of .52 ($p < .001$) between these values. Developers may not be buying in to security culture claims.

7) *How Highly the Team Prioritises Security (SCQ10)*: Security is often treated as an NFR (non-functional requirement) in the software development process, and is not explicitly prioritised by management or teams [53], [59]. Continuing to probe the security culture in the developers' working environment, we asked them how highly they thought their team prioritised software security. See Fig. 7; while 39% felt that their team prioritised it a lot or a great deal, 25% felt that they prioritised it only a little, or not at all. Correlation between answers to this question and the CA score is low to moderate ($r = .48$, $p < .001$).




Fig. 7. SCQ10: How highly do you think your team prioritises software security? (n=896).

8) *How Often Security is Mentioned in Team Communications (SCQ11)*: A developer may perceive their team's prioritisation of software security inaccurately, or their reporting may be subject to social desirability bias. Haney et al. [22] discussed how a security culture involves a 'commitment to address security' and the 'perpetuation of a security mindset.' Communication of priorities is essential within working environments [33], [60]. Furthermore, if indications of concern for secure coding are missing, that in itself allows developers to infer that secure coding is not considered important on the

team [42]. In order to explore security prioritisation further, we asked participants approximately how often software security is mentioned in team communications. Fig. 8 presents the results, which differ markedly from some of the previous answers. 21% of respondents said that security was mentioned on the team about once a week or more often, and 21% said it was mentioned a few times a month. However, a large 58% said it was mentioned about once a month or less, with 9% stating that it is never mentioned. Correlation between answers to this question and the CA score is low to moderate ($r = .44$, $p < .001$).




Fig. 8. SCQ11: How often is software security mentioned in team communications? (n=862).

9) *Time per Week Spent on Secure Coding (SCQ12)*: Time is a scarce resource, and time spent should give further insight into the real emphasis on security in the work environment [42], [61]. We asked participants how much time they spent on security during the week, supplying options from 'None,' 'Less than half an hour,' etc., to 'A week,' and also allowing free text replies. This question can be criticised on the basis that security should come with quality. For example, one respondent replied: 'impossible to quantify this way; all work that leads to higher quality software also usually leads to more secure software.' However, developing secure software entails many security-specific activities; for example, threat modelling is considered an essential component of secure coding [62]–[64]. Security conscious respondents could be expected to have an approximate mental model of how much time they spend on concerns that are primarily security motivated. Therefore, answers to this question are of interest, giving us additional insight into the security culture in the developer's working environment.

Even bearing in mind the qualification above, the answers to this question are not encouraging; Fig. 9 is dominated by red and orange blocks. Less than two hours a week is spent on security by 55% of respondents, with nearly 20% spending none. One free text response was: 'Fluctuates on whether my project is blocked by security review. Usually not at all, sometimes a few hours.' Another participant added 'less than half an hour a month.' Correlation between answers to this question and the CA score is low ($r = .38$, $p < .001$).

C. *Research Question 3: Do secure coding practice and culture correlate, and if not, what lessons can we learn to help support development of secure coding?*

As we examined the answers to our security culture questions we noted the correlation score of each question's answers to




Fig. 9. SCQ12: Roughly how much time do you spend on software security in an average week? (n=878).

the overall CA score for software security practice. In many cases, there was no correlation. Where it did exist it was weak to moderate. By contrast, if good security practice entailed good security culture we would expect to find high correlations between culture answers and practice findings. In an attempt to attain a greater understanding of our results, we broke down some of the answers by CA score. Fig. 10 shows the degree to which the participant thinks their team prioritises security, broken down by CA score. There is clearly a correlation, and for high CA scores we are led to believe that security is very highly prioritised by the team.




Fig. 10. SCQ10: How highly do you think your team prioritises software security? (n=896) This graph shows the value broken down by the number of common security activities undertaken in the environment. The correlation is visible.

However, when we look at how frequently the team communicates about security, a prioritisation of security is not at all clear (see Fig. 11). Even in the teams with the highest security practices, security is mentioned within the team relatively infrequently. At the very highest level of twelve CAs, security is mentioned less than once a week in 36% of teams. Low levels of communication about a topic is an indicator that the topic is low-priority for a team [42]. We suggest that this question gives some insight into environments where, although security is apparently high-priority, the security culture is unfavourable. The minimum work necessary for security compliance is done.

Similarly, when we look at the amount of time spent on security per week it is surprisingly low at all CA levels. At the zero-CA level, 95% of respondents spend less than half an hour a week on security-related activities, with two thirds of these choosing 'None.' Even at the highest level, 59% of participants spend an average of less than two hours per week

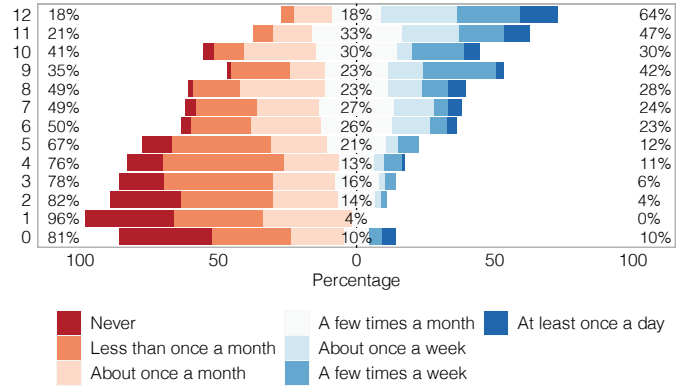


Fig. 11. SCQ11: How often is software security mentioned in team communications? (n=862). This graph shows the value broken down by the number of common security activities undertaken in the environment. Even at the highest CA scores, team security communication is relatively infrequent.

on security activities. Fig. 12 suggests that this question could be useful when attempting to identify a compliance-focused mentality.




Fig. 12. SCQ12: Roughly how much time do you spend on software security in an average week? (n=878). Even at the highest level of common security activities, many participants spend less than two hours a week on security.

VI. DISCUSSION AND CONCLUSION

Asked to comment on the survey, one participant said of the 12 CAs: 'A lot of the questions about bugs or security issues being tracked are kind of missing the point; the system exists, but often times fails. Bugs and bug fixes are in the correct system, but there is not enough time allocated for reporters to verify or validate fixes or for regression.'

This is an excellent illustration of the importance of assessing security culture alongside security activities. If the organisation's employees are going through the motions, producing paperwork and audit trails for compliance purposes, but are not allowed time to do the job properly, a good security culture is missing. In this environment it is inevitable that security issues will slip through the cracks, adversely affecting the

organisation or open source team, impacting their clients, and sometimes posing a systemic threat to their very existence.

Detailed understanding of organisational culture requires longitudinal ethnographic research. In many contexts (such as surveys) this is not achievable, but relative values are of benefit. Our questions are designed to probe core cultural indicators in an organisation. Where values are low, security culture is lacking.

A. Threats to Validity

The questions we used to objectively assess environmental secure coding practice are based on the secure coding activities identified by analysis of BSIMM results as those adopted first and most frequently by security-aware organisations. These activities might not be suited or appropriate to all coding environments. However, the BSIMM data is the largest trove of such data available at this time. It is the closest approach we have identified to finding an empirically-established objective measure of secure coding in environments developing complex applications. The CA Score is a simple measure which does not explore how activities are performed or whether they are implemented correctly. In the necessary research trade-off between cost and detail, it is a lightweight measurement suitable for surveys and quantitative comparisons. For a more thorough measurement of secure software development, more detailed analyses [21] are appropriate.

The questions used to provide insight into participants' understanding of the security culture of their working environment were centred around topics that are prevalent in academic software security literature. Some include ideas that can also be found in organisational climate theory. It is possible that other wordings or additional questions exist which would further enhance our understanding of security culture. Finding and evaluating such questions can be the subject of further research.

As with all questionnaire-based studies, another risk to validity is social desirability bias: the danger that developers will tell us what they think we want to hear. Respondents were self-selecting, and the data cannot be verified. We attempted to mitigate this risk by emphasising in recruitment materials and in the survey consent form that participants did not need to be interested in security and did not need to know anything about it.

Our interpretations of the answers to our questions could be mistaken. For example, the time spent on security by our respondents could in fact be appropriate in an environment with a good security culture. Future research should explore this. In a discussion of assessing software security in the field, it is necessary to bear in mind that the direction of causation cannot be determined. However, correlations and relative emphasis can be observed and can be instructive, which is the approach we have taken here.

B. Conclusion

In this study we add to the existing body of knowledge on how to identify environments where secure coding is prioritised and interest in code security is valued. We wish to ensure that

the correct thing is being measured. In our research, we want to focus on the moon, not on the finger pointing at the moon.

We adopted a list of 12 secure coding activities which have been empirically established as being those adopted first and most often in software security drives. We found that while many of our respondents identified some of these activities in their work environments, in some workplaces none of them were present.

It is well known that security activities can be undertaken for compliance reasons. Researchers into secure coding have established that to achieve secure coding success, a security culture should be established. Based on a thorough literature review, we asked a number of questions designed to evaluate security culture. For example, we asked how much time the respondent spends on secure coding, and how often, on average, secure coding is mentioned in team communications in a week. In this paper we discuss the answers to 12 such questions. We find indications of poor security culture at all levels of security practice. Even where participants state that all 12 common security practices are undertaken in their working environments, communication about security and time spent on security can be very low.

Our results may have an impact in several areas.

Academia: Research on software security practice does not always attempt to measure organisational security [40], and even when this is measured, culture is not quantified. When studying developer behaviour, for example in ethnographic studies, it is important to also consider the security processes and culture in the environment. If this is not done, the research may miss fundamental influences on the developer or team under study. The CA Score provides a simple but empirically validated measure for practice. Our culture questions, especially questions SCQ11 and SCQ12, introduce a way to do a quick litmus test for security culture.

Organisations: For organisations that are serious about encouraging secure coding, it is not enough to introduce practices and follow processes. In addition, time, budget and space for security must be provided.

Industry: Changes to regulation, law and industry standards must look beyond the checkbox approach [25] and consider the holistic background.

The results of this survey show that when evaluating security posture it is not enough merely to measure activities. Security culture must also be evaluated. Understanding the time, money and support available for secure coding activities is crucial to assessing how thoroughly they will be implemented.

ACKNOWLEDGMENTS

We thank the anonymous reviewers, the pilot testers, and all the survey participants for their time and insights. This work is supported by Science Foundation Ireland grants 18/CRT/6222, 13/RC/2077_P2, 13/RC/2094_P2, 15/SIRG/3293.

REFERENCES

- [1] G. McGraw, "Software security," *IEEE Security and Privacy*, vol. 2, no. 5, p. 80–83, 2004.

- [2] I. Ryan, U. Roedig, and K.-J. Stol, "Insecure software on a fragmenting Internet," in *2022 Cyber Research Conference - Ireland (Cyber-RCI)*, 2022.
- [3] K.-J. Stol and B. Fitzgerald, "The ABC of software engineering research," *ACM Transactions on Software Engineering and Methodology*, vol. 27, no. 3, p. 1–51, Sep 2018.
- [4] C. Weir, S. Migueis, M. Ware, and L. Williams, "Infiltrating security into development: exploring the world's largest software security study," in *Proceedings of the 29th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering*. ACM, 2021, p. 1326–1336.
- [5] A. Naiakshina, A. Danilova, C. Tiefenau, M. Herzog, S. Dechand, and M. Smith, "Why do developers get password storage wrong? A qualitative usability study," in *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security*. ACM, 2017, p. 311–328.
- [6] A. Naiakshina, A. Danilova, E. Gerlitz, E. von Zezschwitz, and M. Smith, "If you want, I can store the encrypted password: A password-storage field study with freelance developers," in *Proceedings of the 2019 CHI Conference on Human Factors in Computing Systems*. ACM, May 2019.
- [7] A. Naiakshina, A. Danilova, E. Gerlitz, and M. Smith, "On conducting security developer studies with CS students: Examining a password-storage study with CS students, freelancers, and company developers," in *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems*. ACM, 2020.
- [8] M. G. Jaatun, "Hunting for aardvarks: Can software security be measured?" in *Multidisciplinary Research and Practice for Information Systems*, ser. Lecture Notes in Computer Science, G. Quirchmayr, J. Basl, I. You, L. Xu, and E. Weippl, Eds. Springer, 2012, p. 85–92.
- [9] C. Weir and B. Hermann, "From needs to actions to secure apps? The effect of requirements and developer practices on app security," in *29th USENIX Security Symposium, August 12, 2020 - August 14, 2020*, ser. Proceedings of the 29th USENIX Security Symposium. USENIX Association, 2020, p. 17.
- [10] M. W. Sammy Migueis, John Steven, "BSIMM," <https://www.bsimm.com/>, 2022.
- [11] OWASP, "SAMM," <https://www.opensamm.org/>, 2021, [Online; accessed 13-April-2021].
- [12] Microsoft, "Microsoft Security Development Lifecycle," <https://www.microsoft.com/en-us/securityengineering/sdl/>, 2022, [Online; accessed 01/09/2022].
- [13] SAFECode, "SAFECode Homepage," <https://safecode.org/>, 2022, [Online; accessed 01/09/2022].
- [14] M. G. Jaatun, D. S. Cruzes, K. Bernsmed, I. A. Tøndel, and L. Røstad, "Software security maturity in public organisations," in *Information Security*, ser. Lecture Notes in Computer Science, J. Lopez and C. J. Mitchell, Eds. Springer International Publishing, 2015, p. 120–138.
- [15] P. Morrison, "A security practices evaluation framework," in *Proceedings of the 37th International Conference on Software Engineering - Volume 2*. IEEE Press, 2015, p. 935–938.
- [16] T. D. Oyetoyan, D. S. Cruzes, and M. G. Jaatun, "An empirical study on the relationship between software security skills, usage and training needs in agile settings," in *2016 11th International Conference on Availability, Reliability and Security (ARES)*, 2016, p. 548–555.
- [17] K. Rindell, J. Ruohonen, and S. Hyrynsalmi, "Surveying secure software development practices in Finland," in *Proceedings of the 13th International Conference on Availability, Reliability and Security*, ser. ARES 2018. ACM, 2018, event-place: Hamburg, Germany.
- [18] H. Assal and S. Chiasson, "Security in the software development lifecycle," in *Proceedings of the Fourteenth USENIX Conference on Usable Privacy and Security*. USA: USENIX Association, 2018, p. 281–296.
- [19] D. Votipka, D. Abrokwa, and M. L. Mazurek, "Building and validating a scale for secure software development self-efficacy," in *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems*. ACM, 2020, p. 1–20, event-place: Honolulu, HI, USA.
- [20] P. Morrison, D. Moye, R. Pandita, and L. Williams, "Mapping the field of software life cycle security metrics," *Information and Software Technology*, vol. 102, no. May, p. 146–159, 2018.
- [21] I. A. Tøndel, D. S. Cruzes, M. G. Jaatun, and G. Sindre, "Influencing the security prioritisation in an agile software development project," *Computers and Security*, vol. 118, p. 102744, 2022.
- [22] J. Haney, M. Theofanos, Y. Acar, and S. Spickard Prettyman, "'We make it a big deal in the company': Security mindsets in organizations that develop cryptographic products," in *Proceedings of the Fourteenth Symposium on Usable Privacy and Security*. USENIX Association, 2018, p. 357–373.
- [23] D. Ashenden and D. Lawrence, "Security dialogues: Building better relationships between security and business," *IEEE Security Privacy*, vol. 14, no. 3, pp. 82–87, 2016.
- [24] J. A. Morales, T. P. Scanlon, A. Volkmann, J. Yankel, and H. Yasar, "Security impacts of sub-optimal devsecops implementations in a highly regulated environment," in *Proceedings of the 15th International Conference on Availability, Reliability and Security*. ACM, 2020.
- [25] S. Rahaman, G. Wang, and D. D. Yao, "Security certification in payment card industry: Testbeds, measurements, and recommendations," in *Proceedings of the 2019 ACM SIGSAC Conference on Computer and Communications Security*, ser. CCS '19. New York, NY, USA: ACM, Nov 2019, p. 481–498.
- [26] "PCI DSS," https://www.pcisecuritystandards.org/document_library/?category=pcidss&document=pci_dss, [Online; accessed 26 July 2022].
- [27] C. Heitzentrater and A. Simpson, "A case for the economics of secure software development," in *Proceedings of the 2016 New Security Paradigms Workshop*, ser. NSPW '16. New York, NY, USA: ACM, Sep 2016, p. 92–105.
- [28] I. Rauf, M. Petre, T. Tun, T. Lopez, P. Lunn, D. Van der Linden, J. Towse, H. Sharp, M. Levine, A. Rashid *et al.*, "The case for adaptive security interventions," *ACM Transactions on Software Engineering and Methodology (TOSEM)*, vol. 31, no. 1, pp. 1–52, 2021.
- [29] I. Pashchenko, D.-L. Vu, and F. Massacci, "A qualitative study of dependency management and its security implications," in *Proceedings of the 2020 ACM SIGSAC Conference on Computer and Communications Security*, ser. CCS '20. ACM, Oct 2020, p. 1513–1531.
- [30] I. A. Tøndel, D. S. Cruzes, M. G. Jaatun, and K. Rindell, "The security intention meeting series as a way to increase visibility of software security decisions in agile development projects," in *Proceedings of the 14th International Conference on Availability, Reliability and Security*, ser. ARES '19. New York, NY, USA: ACM, 2019.
- [31] H. Assal and S. Chiasson, "'Think secure from the beginning': A survey with software developers," in *Proceedings of the 2019 CHI Conference on Human Factors in Computing Systems*. ACM, 2019, p. 1–13.
- [32] T. Lopez, H. Sharp, T. Tun, A. K. Bandara, M. Levine, and B. Nuseibeh, "'Hopefully we are mostly secure': Views on secure code in professional practice," in *Proceedings of the 12th International Workshop on Cooperative and Human Aspects of Software Engineering*. IEEE Press, 2019, p. 61–68.
- [33] A. Tuladhar, D. Lende, J. Ligatti, and X. Ou, "An analysis of the role of situated learning in starting a security culture in a software company," in *Seventeenth Symposium on Usable Privacy and Security (SOUPS 2021)*, 2021, p. 617–632.
- [34] M. Tahaei, A. Frik, and K. Vaniea, "Privacy champions in software teams: Understanding their motivations, strategies, and challenges," in *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems*. New York, NY, USA: ACM, 2021, p. 16.
- [35] I. Ryan, U. Roedig, and K.-J. Stol, "Understanding developer security archetypes," in *International Workshop on Engineering and Cybersecurity of Critical Systems (EnCyCriS)*. Association of Computer Machinery, Aug. 2021.
- [36] J. M. Haney and W. G. Lutters, "'It's scary... it's confusing... it's dull': How cybersecurity advocates overcome negative perceptions of security," in *Fourteenth Symposium on Usable Privacy and Security (SOUPS 2018)*, 2018, p. 411–425.
- [37] C. Weir, I. Becker, and L. Blair, "A passion for security: Intervening to help software developers," in *2021 IEEE/ACM 43rd International Conference on Software Engineering: Software Engineering in Practice (ICSE-SEIP)*, 2021, p. 21–30.
- [38] S. Xiao, J. Witschey, and E. Murphy-Hill, "Social influences on secure development tool adoption: Why security tools spread," in *Proceedings of the 17th ACM Conference on Computer Supported Cooperative Work and Social Computing*. New York, NY, USA: ACM, 2014, p. 1095–1106.
- [39] J. Witschey, S. Xiao, and E. Murphy-Hill, "Technical and personal factors influencing developers' adoption of security tools," in *ACM Workshop on Security Information Workers*. ACM Press, 2014, p. 23–26.
- [40] H. Palombo, A. Z. Tabari, D. Lende, J. Ligatti, and X. Ou, "An ethnographic understanding of software (in)security and a co-creation model to improve secure software development," in *Sixteenth Symposium on Usable Privacy and Security (SOUPS 2020)*, 2020, pp. 205–220.
- [41] B. Schneider, M. G. Ehrhart, and W. H. Macey, "Organizational climate and culture," *Annual Review of Psychology*, vol. 64, no. 1, p. 361–388, 2013.

- [42] R. Arizon-Peretz, I. Hadar, G. Luria, and S. Sherman, "Understanding developers privacy and security mindsets via climate theory," *Empirical Software Engineering*, vol. 26, no. 6, p. 34, 2021.
- [43] R. Arizon-Peretz, I. Hadar, and G. Luria, "The importance of security is in the eye of the beholder: Cultural, organizational, and personal factors affecting the implementation of security by design," *IEEE Transactions on Software Engineering*, vol. 48, p. 4433–4446, 2022.
- [44] I. Ryan, U. Roedig, and K.-J. Stol, "Measuring secure coding practice and culture: A finger pointing at the moon is not the moon," Feb 2023. [Online]. Available: osf.io/mz29b
- [45] W. Foddy and W. H. Foddy, *Constructing Questions for Interviews and Questionnaires: Theory and Practice in Social Research*. Cambridge University Press, 1994.
- [46] J. Witschey, O. Zielinska, A. Welk, E. Murphy-Hill, C. Mayhorn, and T. Zimmermann, "Quantifying developers' adoption of security tools," in *2015 10th Joint Meeting of the European Software Engineering Conference and the ACM SIGSOFT Symposium on the Foundations of Software Engineering, ESEC/FSE 2015 - Proceedings*, 2015, p. 260–271.
- [47] A. Danilova, A. Naiakshina, and M. Smith, "One size does not fit all: a grounded theory and online survey study of developer preferences for security warning types," in *Proceedings of the ACM/IEEE 42nd International Conference on Software Engineering*. ACM, 2020, p. 136–148.
- [48] A. Danilova, A. Naiakshina, S. Horstmann, and M. Smith, "Do you really code? Designing and evaluating screening questions for online surveys with programmers," in *2021 IEEE/ACM 43rd International Conference on Software Engineering (ICSE)*, May 2021, p. 537–548.
- [49] H. Kaur, S. Amft, D. Votipka, Y. Acar, and S. Fahl, "Where to recruit for security development studies from: Comparing six software developer samples," in *31st USENIX Security Symposium (USENIX Security 22)*, Aug 2022, p. 24.
- [50] E. Murphy-Hill, C. Jaspan, C. Sadowski, D. Shepherd, M. Phillips, C. Winter, A. Knight, E. Smith, and M. Jorde, "What predicts software developers' productivity?" *IEEE Transactions on Software Engineering*, vol. 47, no. 3, pp. 582–594, 2019.
- [51] R Core Team, *R: A Language and Environment for Statistical Computing*, R Foundation for Statistical Computing, Vienna, Austria, 2022. [Online]. Available: <https://www.R-project.org/>
- [52] G. Norman, "Likert scales, levels of measurement and the "laws" of statistics," *Advances in Health Sciences Education*, vol. 15, no. 5, p. 625–632, Dec 2010.
- [53] J. Xie, H. R. Lipford, and B. Chu, "Why do programmers make security errors?" in *2011 IEEE Symposium on Visual Languages and Human-Centric Computing (VL/HCC)*. IEEE, 2011, p. 161–164.
- [54] O. Pieczul, S. Foley, and M. E. Zurko, "Developer-centered security and the symmetry of ignorance," in *Proceedings of the 2017 New Security Paradigms Workshop*, 2017, pp. 46–56.
- [55] T. W. Thomas, M. Tabassum, B. Chu, and H. Lipford, "Security during application development: An application security expert perspective," in *Proceedings of the 2018 CHI Conference on Human Factors in Computing Systems*. New York, NY, USA: ACM, 2018.
- [56] I. Rauf, D. van der Linden, M. Levine, J. Towse, B. Nuseibeh, and A. Rashid, "Security but not for security's sake: The impact of social considerations on app developers' choices," in *Proceedings of the IEEE/ACM 42nd International Conference on Software Engineering Workshops*, ser. ICSEW'20. ACM, 2020, p. 141–144, event-place: Seoul, Republic of Korea.
- [57] M. Tahaei, A. Jenkins, K. Vaniea, and M. Wolters, "'I don't know too much about it': On the security mindsets of computer science students," in *Proceedings of the 2019 ACM Conference on Innovation and Technology in Computer Science Education*, 2019, pp. 350–350.
- [58] N. Tomas, J. Li, and H. Huang, "An empirical study on culture, automation, measurement, and sharing of DevSecOps," in *2019 International Conference on Cyber Security and Protection of Digital Services (Cyber Security)*, Jun 2019, p. 1–8.
- [59] M. Tahaei and K. Vaniea, "A survey on developer-centred security," in *2019 IEEE European Symposium on Security and Privacy Workshops (EuroS&PW)*. IEEE, 2019, pp. 129–138.
- [60] A. Poller, L. Kocksch, K. Kinder-Kurlanda, and F. A. Epp, "First-time security audits as a turning point?: Challenges for security practices in an industry software development team," in *Proceedings of the 2016 CHI Conference Extended Abstracts on Human Factors in Computing Systems*, ser. CHI EA '16. ACM, 2016, p. 1288–1294, event-place: San Jose, California, USA.
- [61] D. Oliveira, M. Rosenthal, N. Morin, K.-C. Yeh, J. Cappos, and Y. Zhuang, "It's the psychology stupid: how heuristics explain software vulnerabilities and how priming can illuminate developer's blind spots," in *Proceedings of the 30th Annual Computer Security Applications Conference*. ACM, 2014, p. 296–305.
- [62] Y. Acar, C. Stransky, D. Wermke, C. Weir, M. L. Mazurek, and S. Fahl, "Developers need support, too: A survey of security advice for software developers," in *2017 IEEE Cybersecurity Development (SecDev)*. IEEE, 2017, p. 22–26.
- [63] C. Weir, I. Becker, J. Noble, L. Blair, M. A. Sasse, and A. Rashid, "Interventions for long-term software security: Creating a lightweight program of assurance techniques for developers," *Software - Practice and Experience*, vol. 50, no. 3, p. 275–298, 2020.
- [64] J. Whitmore and W. Tobin, "Improving attention to security in software design with analytics and cognitive techniques," in *2017 IEEE Cybersecurity Development (SecDev)*, 2017, p. 16–21.